



## Assignment in Optimization 2018 - Quasi-Newton methods

John Rapp Farnes, Erik Ris

Axel Sjöberg

971126-4771, 940309-8677

960314-3653

jo5113fa-s@student.lu.se, ine15erik@student.lu.se

ax3817sj-s@student.lu.se

January 30, 2019

# Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
<b>2</b>	<b>Method</b>	<b>1</b>
2.1	linesearch . . . . .	1
2.2	nonlinearmin . . . . .	2
<b>3</b>	<b>Optimization problems</b>	<b>2</b>
<b>4</b>	<b>Results</b>	<b>3</b>
4.1	Optimization problems . . . . .	3
4.2	Example of printouts . . . . .	3
4.3	Comparison of different initial point and methods . . . . .	4
<b>5</b>	<b>Discussion</b>	<b>4</b>
5.1	Line search . . . . .	4
5.1.1	Choice of line search algorithm . . . . .	4
5.1.2	Numerically calculating $F'(0)$ . . . . .	5
5.2	Termination criterion . . . . .	5
5.3	Performance on problems (1) - (4) . . . . .	6
5.3.1	Rosenbrock (1) . . . . .	6
5.3.2	Penalty function problems (2), (3) . . . . .	6
5.3.3	Barrier Function Problem (4) . . . . .	7
5.4	Comparison of DFP and BFGS . . . . .	7
	<b>Appendix - Matlab implementation</b>	<b>8</b>
	linesearch.m . . . . .	8
	nonlinearmin.m . . . . .	10
	penalty.m . . . . .	12
	barrier.m . . . . .	12

# 1 Introduction

The purpose of this project was to implement and compare two different kinds of Quasi-Newton methods: BFGS (Broyden–Fletcher–Goldfarb–Shanno) and DFP (Davidon–Fletcher–Powell) for minimizing functions. The different methods are used to solve different optimization problems, including constrained optimization via penalty- and barrier-functions. The project was based upon the course FMAN60 and the literature used in the course.

## 2 Method

The basic idea behind both optimization algorithms is to perform iterations of the following routine:

1. Start at a point  $x_k$ , where  $x_0$  is a given initial point
2. Determine a search direction  $d_k$
3. Minimize  $F(\lambda) = f(x_k + \lambda d_k)$  along  $d_k$  using a line search algorithm. Let  $\lambda^* = \text{optimum}$
4. Let  $x_{k+1} = x_k + \lambda^* d_k$
5. Check  $x_{k+1}$  against the termination criterion

This algorithm was implemented in MatLab as *nonlinearmin*, with *linesearch* as a subroutine, with the following function signatures:

$$\begin{aligned} [\lambda, No\_of\_iterations] &= linesearch(func, x, d) \\ [x, No\_of\_iterations] &= nonlinearmin(f, start, method, tol, printout) \end{aligned}$$

### 2.1 linesearch

The Golden Section method was first implemented as the line search method but it was later discarded in favour of Armijo's rule. Armijo's rule is an inexact line search method that specifies two criteria that a  $\lambda$  solving the line search optimization problem should satisfy, based on two parameters  $\alpha > 1$  and  $0 < \varepsilon < 1$ . The algorithm starts with an initial  $\lambda = \lambda_0$  and iteratively checks the  $\lambda$  against the criteria of Armijo's rule. If the criteria that the  $\lambda$  is too large, it is divided by  $\alpha$ , and multiplied by  $\alpha$  if it is too small. The criteria are based on "overshooting" the tangent at  $\lambda = 0$ . For a further explanation of Armijo's rule see the course literature, and the schematic overview in figure 1.

The chosen parameters in the implementation were  $\alpha = 2$ ,  $\varepsilon = 0.1$  and  $\lambda_0 = 1$ .




Figure 1: Schematic overview of Armijo's rule, where  $T$  is the "overshooting" tangent at  $\lambda = 0$

## 2.2 nonlinearmin

Both of the examined methods determine the direction  $d_k$  using the gradient of  $f$  as well as conjugate directions. The difference between the methods concerned the formula for calculating the matrix  $D$ , later used for determining  $d_k = -D \cdot \nabla f$ . The methods were implemented according to the course literature.

## 3 Optimization problems

The function *nonlinearmin* was used to solve four different optimization problems, referred to as (1), (2), (3), (4) below.

Problem (2) and (3) were solved using the penalty method which was implemented in MatLab, using the sequence  $\mu = 0.01, 0.1, 1, 10, 100$ . Problem (4) was solved with the barrier method with  $\varepsilon = 1, 0.1, 0.01, 0.001, 0.0001$ .

$$\begin{aligned} \underset{x}{\text{minimize}} \quad & f(\bar{x}) = 100(x_2 - x_1^2)^2 + (1 - x_1)^2 \\ & \text{(Rosenbrock's function)} \end{aligned} \tag{1}$$

$$\begin{aligned} \underset{x}{\text{minimize}} \quad & f(\bar{x}) = e^{x_1 x_2 x_3 x_4 x_5} \\ \text{subject to} \quad & x_1^2 + x_2^2 + x_3^2 + x_4^2 + x_5^2 = 10 \\ & x_2 x_3 = 5 x_4 x_5 \\ & x_1^3 + x_3^3 = -1 \end{aligned} \tag{2}$$

$$\begin{aligned} \underset{x}{\text{minimize}} \quad & f(\bar{x}) = e^{x_1} + x_1^2 + x_1 x_2 \\ \text{subject to} \quad & \frac{1}{2} x_1 + x_2 - 1 = 0 \end{aligned} \tag{3}$$

$$\begin{aligned} \underset{x}{\text{minimize}} \quad & f(\bar{x}) = (x_1 - 5)^2 + (x_2 - 3)^2 \\ \text{subject to} \quad & x_1 + x_2 \leq 3 \\ & -x_1 + 2x_2 \leq 4 \end{aligned} \tag{4}$$

## 4 Results

### 4.1 Optimization problems

Table 1 and 2 shows the results from the numerical optimization of the four problems, from the starting points given in the assignment.

Problem	Optimal point $x^*$	$f(x^*)$	Iterations
Rosenbrock	(0.9974, 0.9947)	$6.9089 \cdot 10^{-6}$	37
Penalty problem	(-1.7170, 1.8275, 1.5956, -0.7637, -0.7637)	0.0539	13
Exercise 9.3	(-1.2785, 1.6393)	-0.1828	17
Exercise 9.5	(2.4977, 0.4979)	12.5224	19

Table 1: Results from *nonlinearmin* using DFP

Problem	Optimal point $x^*$	$f(x^*)$	Iterations
Rosenbrock	(0.9974, 0.9948)	$6.6996 \cdot 10^{-6}$	37
Penalty problem	(-1.7165, 1.8284, 1.5950, -0.7637, -0.7637)	0.0539	13
Exercise 9.3	(-1.2785, 1.6393)	-0.1828	14
Exercise 9.5	(2.4975, 0.4980)	12.5224	19

Table 2: Results from *nonlinearmin* using BFGS

### 4.2 Example of printouts

iterations	x	stepsize	f(x)	norm(grad)	ls	iters	lambda
1	0.2694 0.0473	0.2735	0.5974	3.2717	8	0.0078	
2	0.4168 0.1498	0.1795	0.3973	1.3007	4	0.1250	
3	0.5836 0.3153	0.2350	0.2371	1.9662	3	0.2500	
4	0.7581 0.5429	0.2867	0.1596	0.7908	2	0.5000	
5	0.8869 0.7688	0.2601	0.0442	3.0943	3	1.0000	
6	0.9306 0.8635	0.1043	0.0054	4.2719	3	1.0000	
7	0.9934 0.9811	0.1333	0.0033	0.6970	4	2.0000	
8	1.0054 1.0113	0.0325	0.0000	0.0831	5	4.0000	
9	0.9968 0.9934	0.0198	0.0000	0.0056	5	4.0000	
10	0.9967 0.9934	0.0000	0.0000	0.0412	11	0.0010	

Figure 2: Print out from *nonlinearmin.m*,  $f = \text{rosenbrock.m}$ , method = "BFGS"

iterations	x	stepsize	f(x)	norm(grad)	ls iters	lambda
1	-1.7914 1.8017 1.6920 -0.8660 -0.8660	0.4622	0.0283	0.0127	5	4.0000
2	-1.7845 1.8216 1.6795 -0.8614 -0.8614	0.0253	0.0282	0.0124	4	2.0000
3	-1.7556 1.8728 1.6487 -0.8632 -0.8632	0.0664	0.0281	0.0060	7	16.0000
4	-1.7575 1.8666 1.6534 -0.8626 -0.8626	0.0081	0.0281	0.0009	4	2.0000
5	-1.7583 1.8665 1.6530 -0.8629 -0.8629	0.0011	0.0281	0.0021	2	0.5000

Figure 3: Print out from nonlinearmin.m,  $f = (2) + \text{penalty func.}$ , method = "DFP",  $\mu = 0.01$

### 4.3 Comparison of different initial point and methods

Method	Intial point	Optimal point $x^*$	$f(x^*)$	Iterations
DFP	(0, 2, 2, -1, -1)	(-1.7171, 1.8274, 1.5956, -0.7637, -0.7637)	0.0539	13
DFP	(-1, 2, 2, -1, -1)	(-1.7170, 1.8275, 1.5956, -0.7637, -0.7637)	0.0539	15
BFGS	(0, 2, 2, -1, -1)	(-1.7180, 1.8256, 1.5967, -0.7636, -0.7636)	0.0539	12
BFGS	(-2, 2, 2, -1, -1)	(-1.7165, 1.8284, 1.5950, -0.7637, -0.7637)	0.0539	13

Table 3: Comparison of results from *nonlinearmin* for problem (2) with different initial points and methods

## 5 Discussion

### 5.1 Line search

#### 5.1.1 Choice of line search algorithm

At first, the Golden Section method was chosen as the line search method because of its reliable convergence properties. The problem with the method however was that in order to find optimum located far from the initial point, the initial interval had to be very large. This led to computations that did not finish in a reasonable amount of time. In order to obtain flexibility regarding the location of the optimum, while not sacrificing performance, Armijo's rule was chosen instead.

### 5.1.2 Numerically calculating $F'(0)$

Armijo's rule requires the derivative to be computed at  $\lambda = 0$ . This derivative must be calculated numerically, approximating the limit in the definition of the derivative, in order to achieve reasonable performance. The problem with this approach is to determine which  $h$  to use in the limit. For some functions, such as the *test\_func* in the assignment, a  $h$  that is very small ( $\approx 10^{-60}$ ) must be used in order to avoid numerical errors (which will result in an *Inf* or *NaN* derivative). The problem with such a small  $h$  is that it numerically "kills" functions that are less steep, resulting in a derivative of 0. As such,  $h$  must be chosen dynamically depending on the local behavior of the function.

At first, this problem was solved by starting with a small  $h$  and increasing it as long as the resulting derivative was 0. This approach however proved to be either slow (if too many different values of  $h$  were tried) or unstable (if too few values were tried). After consulting other groups having the same problem, we implemented a solution where first a  $\lambda = \lambda'$  is found that has the property that  $F(\lambda') \approx F(0)$ . This  $\lambda'$  was found by starting with  $\lambda = \lambda_0$  and dividing  $\lambda$  with  $\alpha$  as long as  $F(\lambda) > F(0)$ , or multiplying it with  $\alpha$  as long as  $F(\lambda) < F(0)$ . The resulting  $\lambda'$  was then used in calculating  $h$ , with  $h = h_0\lambda'$  where  $h_0$  is a constant ( $h_0 = 10^{-6}$  in the implementation). This gives a good value of  $h$  since  $\lambda'$  is a measure of the steepness of the function at  $\lambda = 0$ . If it is steep, the  $\lambda'$  is small resulting in a small  $h$  (required for calculating the derivative) and the opposite if the function is flat (not "killing" the derivative). This method is implemented as *find\_suitable\_lambda*, which returns  $\lambda'$ . This  $\lambda'$  is then used as  $\lambda_0$  in Armijo's rule.

In addition, we added a condition where we set the derivative to 0 if it is calculated to be a positive number, in order to terminate earlier from Armijo's rule.

## 5.2 Termination criterion

For the termination criterion, we combined criteria regarding the changes in  $\bar{x}_k$  and  $f(\bar{x})$  using (5) and (6), where  $\epsilon$  is the tolerance.

$$|\bar{x}_{k+1} - \bar{x}_k| \leq \epsilon \quad (5)$$

$$|f(\bar{x}_{k+1}) - f(\bar{x}_k)| < \epsilon \quad (6)$$

The algorithm was terminated if any of the above conditions were met. The motivation behind doing this is that the criterion then considers both when changes in  $\bar{x}$  is declining and  $f(\bar{x})$  is. The inequalities and equalities in (5) and (6) have been taken from the course literature.

This termination criterion was not only applied to outer loop of the *nonlinearmin* but also the inner loop, as changes that were too small during the inner iterations may result in elements of  $D$  that were *NaN*, since  $q = 0$  if  $y_{k+1} - y_k$  was too small. In addition, the algorithm terminated if  $\nabla f < \epsilon$  as close to zero gradients may have resulted in zero directions and thereby line search optimization problems that had no solution.

### 5.3 Performance on problems (1) - (4)

#### 5.3.1 Rosenbrock (1)

Depending on the initial point, sometimes the DFP method outperformed the BFGS method and sometimes the BFGS outperformed the DFP. For both the DFP- and the BFGS method, more iterations were required as the initial point was further away from the optimal minimum point. Choosing an initial point that was very far away from the optimal point sometimes meant that the problem converged to a local minimum that was not the global minimum and sometimes lead to errors in the line search. This was the case for both the DFP- and BFGS method. For instance setting the initial point to  $\bar{x} = (10000, 50000)$  and using the DFP method, the problem converged to the optimal point  $\bar{x} = (-1100, 1.3282)$ , which gave the minimum value  $1.3305 \cdot 10^6$ . When testing initial points that were close to the global min, the linesearch always worked and always converged towards the same optimal solution. In order to make the algorithm converge to the global minimum, several different initial points that are far away from each other should be tested. If only one initial point is tested, the algorithm could converge towards a local minimum that is not a global minimum, which was the case above.

#### 5.3.2 Penalty function problems (2), (3)

The minimization of (2) gave the same minimal value for DFP and BFGS, however the minimal value differed depending on the initial points. Depending on the initial point the BFGS algorithm and DFP algorithm required a different amount of iterations before terminating the optimization due to the termination criterion. For all the tested initial points the BFGS method required less than or the equal amount iterations as the DFP method. These points were however arbitrarily chosen and this may not hold in general.

For minimization problem (2) there was no apparent relation between the required number of iterations and the distance from the initial point and the optimal point, as can be seen in table 3 above. For both the DFP method and BFGS method there existed an initial point with greater distance to the optimal point than another initial point, at the same time as it required fewer iterations before terminating the optimization.

The minimization problem (3) converged to the same solution for all the tested initial points, except for initial points that were far away from the optimal point. The initial points far away from the optimal solution generated an error in the linesearch as it required too many iterations in the linesearch. As with minimization problem (2), depending on the initial point, the BFGS algorithm and DFP algorithm required a different amount of iterations before terminating the optimization. Setting different initial points, the BFGS method always required less than or the equal amount of iterations as the DFP method required. None the less, the initial points were still chosen arbitrarily as above and analogously we cannot determine if the BFGS method requires less than or the equal amount of iterations as the DFP method for all initial points.

Choosing an initial point that is very far away from the optimal point resulted in an error in the linesearch as it required too many iterations in the linesearch. Setting an initial point somewhat close to the optimal value however always meant that the problem converged to same solution for (2) and (3). Choosing a suitable initial point proved to be quite the difficult task. No one size fits all method that works in all cases were apparent. The most successful strategy was



choosing plenty of initial points that were far away from each other and compare the solutions to each other. A few different points close to each other were tried out to test the consistency of the program behaviour.

### 5.3.3 Barrier Function Problem (4)

For the initial point  $\bar{x} = (0, 0)$  the BFGS and DFP converged to the same solution and the optimization was terminated in the same number of iterations. Changing the initial point to  $\bar{x} = (0, -10)$  gave the same minimum value, however it only required 14 iterations for the DFP method and 17 for the BFGS method. For all feasible points, the DFP required less than or the equal amount of iterations as the BFGS method. As with minimization problem (2) and (3) the initial points were chosen arbitrarily and we cannot determine whether or not the DFP method for all initial points require less or the equal amount of iterations as the BFGS method, in problem (4).

For all the feasible points tested, the barrier function converged to the same solution. The strategy concerning choosing a suitable initial point in this problem is thereby reduced to finding feasible points.

## 5.4 Comparison of DFP and BFGS

For penalty problems as problem (2) and (3) the BFGS method performed better than the DFP method in the sense that it required less iterations for the same minimal value. For the barrier problem (4) the DFP performed better than the BFGS in the same regard. As such, we are not able to draw a clear conclusion regarding which if DFP or BFGS is "better". As only two barrier functions with a few different initial points and one barrier function with a few different initial points were tested, there is a lack of foundation to build strong conclusions. For the Rosenbrock function depending on the initial point whichever of the DFP and BFGS method outperformed the other method varied.

## Appendix - Matlab implementation

### linesearch.m

```
1 function [lambda, No_of_iterations] = linesearch(func, x, d)
2
3     F = @(lambda) func(x + lambda .* d);
4
5     [lambda, No_of_iterations] = armijo(F);
6
7     if isnan(func(x+lambda*d))
8         error('Line search is NaN!')
9     elseif func(x+lambda*d) > func(x)
10        error('Line search did no good!')
11    end
12 end
13
14 function [lambda, No_of_iterations] = armijo(F)
15
16     alpha = 2;
17     epsilon = 0.1;
18     lambda_0 = 1;
19
20     max_iterations = 500;
21
22     F_0 = F(0);
23     [lambda, No_of_iterations] = find_suitable_lambda(F, lambda_0, F_0, alpha, ...
24         max_iterations);
25
26     derivative_0 = derivative(F, 0, lambda);
27
28     T = @(lambda) F_0 + epsilon*lambda*derivative_0;
29
30     while No_of_iterations < max_iterations
31         No_of_iterations = No_of_iterations + 1;
32
33         below_tangent = F(lambda) ≤ T(lambda);
34         too_small = F(alpha*lambda) ≤ T(alpha*lambda);
35
36         if below_tangent && not(too_small)
37             break;
38         end
39
40         if not(below_tangent)
41             lambda = lambda / alpha;
42         elseif too_small
43             lambda = lambda * alpha;
44         end
45     end
46
47     if No_of_iterations == max_iterations
48         error('Too long time in line search!')
49     end
50 end
51
52 function [lambda, No_of_iterations] = find_suitable_lambda(F, lambda, F_0, ...
53     alpha, max_iterations)
54     No_of_iterations = 0;
```

```

54     was_bigger = F(lambda) > F_0;
55
56     while (isnan(F(lambda)) || F(lambda) > F_0) && No_of_iterations < ...
        max_iterations
57         lambda = lambda / alpha;
58         No_of_iterations = No_of_iterations + 1;
59     end
60
61     if was_bigger
62         return;
63     end
64
65     while F(lambda) ≤ F_0 && No_of_iterations < max_iterations
66         lambda = lambda * alpha;
67         No_of_iterations = No_of_iterations + 1;
68     end
69 end
70
71 function dy = derivative(F, x, suitable_lambda)
72     h_0 = 1e-6;
73     h = suitable_lambda * h_0;
74
75     dy = (F(x + h) - F(x - h)) / (2*h);
76
77     if (dy > 0)
78         dy = 0;
79     end
80 end

```

## nonlinearmin.m

```
1 function [x, No.of.iterations] = nonlinearmin(f, start, method, tol, printout)
2
3     update_matrix = update_matrix_fn(method);
4
5     max_iterations = 500;
6     n = numel(start);
7     x = start;
8     No.of.iterations = 0;
9
10    print_header(printout);
11
12    while No.of.iterations < max_iterations
13        No.of.iterations = No.of.iterations + 1;
14
15        y = x;
16        D = eye(n);
17        for j = 1:n
18            gf = grad(f, y);
19
20            if grad_too_small(gf, tol)
21                x = y;
22                return;
23            end
24
25            d = -D * gf;
26            [lambda, ls_iter] = linesearch(f, y, d);
27            next_y = y + lambda*d;
28
29            if should_stop(f, y, next_y, tol)
30                break;
31            end
32
33            D = update_matrix(f, D, lambda, d, y, next_y);
34            y = next_y;
35        end
36        x_next = next_y;
37
38        print_row(printout, No.of.iterations, x, x_next, f, gf, ls_iter, lambda);
39
40        if should_stop(f, x, x_next, tol)
41            break;
42        end
43
44        x = x_next;
45        if No.of.iterations == max_iterations
46            error('Too long time in minimization!')
47        end
48    end
49
50    x = x_next;
51    if isnan(f(x)) || f(x) > f(start)
52        error('Bad job of the unlinear search!')
53    end
54
55 end
56
57 function update_matrix = update_matrix_fn(method)
58     if method == "DFP"
```

```

59         update_matrix = @update_DFP;
60     elseif method == "BFGS"
61         update_matrix = @update_BFGS;
62     end
63 end
64
65 function [p, q] = get_pq(f, lambda, d, y, next_y)
66     p = lambda * d;
67     q = grad(f, next_y) - grad(f, y);
68
69     if all(p == 0)
70         error("p == 0");
71     end
72
73     if all(q == 0)
74         error("q == 0");
75     end
76 end
77 function D = update_DFP(f, D, lambda, d, y, next_y)
78     [p, q] = get_pq(f, lambda, d, y, next_y);
79     D = D + (p*p')/(p'*q) - (D*(q*q')*D)/(q'*D*q);
80 end
81 function D = update_BFGS(f, D, lambda, d, y, next_y)
82     [p, q] = get_pq(f, lambda, d, y, next_y);
83     D = D + (1 + q'*D*q/(p'*q))*(p*p')/(p'*q) - (p*q'*D+D*q*p')/(p'*q);
84 end
85 function stop = grad_too_small(gf, tol)
86     stop = norm(gf) < tol;
87     if stop
88         disp("Gradient too small!");
89     end
90 end
91 function stop = should_stop(f, x, x_next, tol)
92     stop = all(abs(x_next - x) <= tol) || abs(f(x_next) - f(x)) < tol;
93 end
94
95 function print_header(printout)
96     if (printout)
97         disp("iterations      x      stepsize      f(x)      norm(grad) ...
98             ls_iters      lambda");
99     end
100 end
101 function print_row(printout, No_of_iterations, x, x_next, f, gf, ls_iter, lambda)
102     if (printout)
103         step_size = norm(x_next - x);
104         fprintf('%5.0f %12.4f %12.4f %15.4f %13.4f %8.0f %13.4f\n', ...
105             No_of_iterations, ...
106             x_next(1), ...
107             step_size, ...
108             f(x_next), ...
109             norm(gf), ...
110             ls_iter, ...
111             lambda ...
112         );
113         for k = 2:numel(x_next)
114             fprintf('%18.4f\n', x_next(k));
115         end
116     end
117 end

```

## penalty.m

```
1 function penalty = penalty(f, g_k, h_k, mu)
2
3     function q = q(x)
4         q = f(x) + mu * alpha(g_k, h_k, x);
5     end
6     penalty = @q;
7 end
8
9 function val = values(g_k, x)
10    val(1) = 0;
11    for i = 1:numel(g_k)
12        g = g_k{i};
13        val(i) = g(x);
14    end
15 end
16
17 function alpha = alpha(g_k, h_k, x)
18    alpha_g = max(values(g_k, x), 0).^2;
19    alpha_h = values(h_k, x).^2;
20
21    alpha = sum(alpha_g) + sum(alpha_h);
22 end
```

## barrier.m

```
1 function barrier = barrier(f, g_k, epsilon)
2
3     function q = q(x)
4         q = f(x) + epsilon * beta(g_k, x);
5     end
6
7     barrier = @q;
8 end
9
10 function val = values(g_k, x)
11    val(1) = 0;
12    for i = 1:numel(g_k)
13        g = g_k{i};
14        val(i) = g(x);
15    end
16 end
17
18 function beta = beta(g_k, x)
19    val = values(g_k, x);
20
21    not_feasible = any(val ≥ 0);
22
23    if not_feasible
24        beta = 1e200;
25    else
26        val = 1./val;
27        beta = -sum(val);
28    end
29 end
```